

```
# Import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pickle

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, roc_auc_score, roc_curve

# Set plot styles
sns.set(style="whitegrid")
plt.rcParams["figure.figsize"] = (10, 6)

# -----
# Step 1: Load the Dataset
# -----
df = pd.read_csv("/content/dataset.csv")
print("Shape:", df.shape)
print(df.head())
print(df.info())
```

↻ Shape: (349, 10)

	Disease	Fever	Cough	Fatigue	Difficulty Breathing	Age	Gender
0	Influenza	Yes	No	Yes	Yes	19	Female
1	Common Cold	No	Yes	Yes	No	25	Female
2	Eczema	No	Yes	Yes	No	25	Female
3	Asthma	Yes	Yes	No	Yes	25	Male
4	Asthma	Yes	Yes	No	Yes	25	Male

	Blood Pressure	Cholesterol Level	Outcome Variable
0	Low	Normal	Positive
1	Normal	Normal	Negative
2	Normal	Normal	Negative
3	Normal	Normal	Positive
4	Normal	Normal	Positive

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 349 entries, 0 to 348

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	Disease	349 non-null	object
1	Fever	349 non-null	object
2	Cough	349 non-null	object
3	Fatigue	349 non-null	object
4	Difficulty Breathing	349 non-null	object
5	Age	349 non-null	int64
6	Gender	349 non-null	object
7	Blood Pressure	349 non-null	object
8	Cholesterol Level	349 non-null	object
9	Outcome Variable	349 non-null	object

dtypes: int64(1), object(9)

memory usage: 27.4+ KB

None

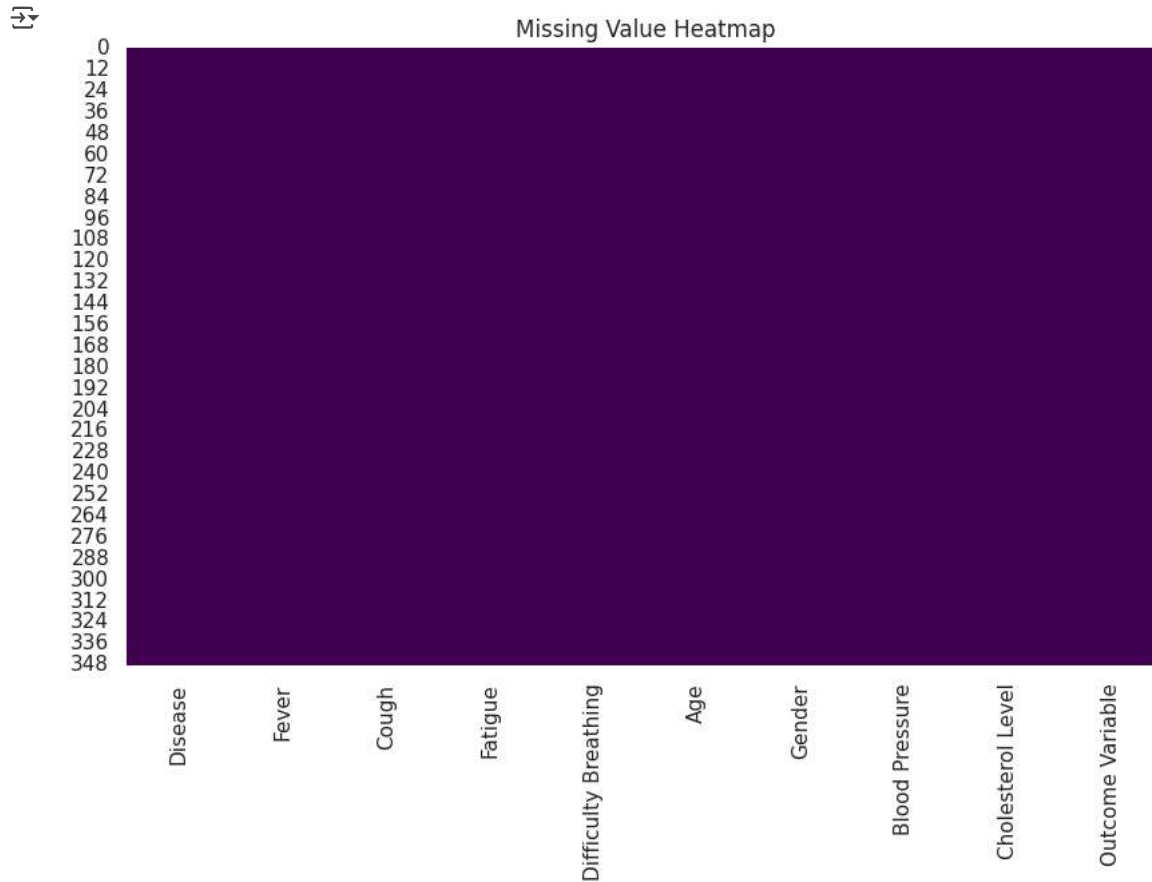
```
# Step 2: Data Cleaning
# -----
# Check for missing values
print("\nMissing values:\n", df.isnull().sum())
```

↻ Missing values:

Disease	0
Fever	0
Cough	0
Fatigue	0
Difficulty Breathing	0
Age	0
Gender	0
Blood Pressure	0
Cholesterol Level	0
Outcome Variable	0

dtype: int64


```
# Visualize missing values
sns.heatmap(df.isnull(), cbar=False, cmap='viridis')
plt.title("Missing Value Heatmap")
plt.show()
```



```
# Drop rows with missing values
df = df.dropna()
```

```
# Drop duplicates
df = df.drop_duplicates()
```


```
# Step 3: Exploratory Data Analysis (EDA)
# -----
# Distribution of target
if 'Disease' in df.columns:
    sns.countplot(x='Disease', data=df)
    plt.title("Target Class Distribution")
    plt.show()
```



```
# Histograms for numeric features
df.select_dtypes(include=[np.number]).hist(bins=30, figsize=(15, 10), color='skyblue')
plt.suptitle("Feature Distributions")
plt.show()
```



Feature Distributions



```
# Import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt # This line imports the necessary library
import seaborn as sns
import pickle


# Load the dataset (Make sure the path is correct)
df = pd.read_csv("/content/dataset.csv")

# ... (rest of the code) ...

# Correlation heatmap (numeric only to avoid string conversion errors)
plt.figure(figsize=(12, 8)) # Now 'plt' is defined and can be used
numeric_df = df.select_dtypes(include=[np.number])
sns.heatmap(numeric_df.corr(), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Matrix (Numeric Features)")
plt.show()
```



Correlation Matrix (Numeric Features)



```
# Step 4: Preprocessing
# -----
from sklearn.preprocessing import LabelEncoder # Import LabelEncoder here
label_encoders = {}
for col in df.select_dtypes(include='object').columns:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le

# Import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pickle

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler # Import StandardScaler here
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, roc_auc_score, roc_curve

if 'Disease' not in df.columns:
    raise ValueError("Target column not found in dataset. Please ensure it's named 'Disease'.")

X = df.drop('Disease', axis=1)
y = df['Disease']

# Scale features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split data
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

# Step 5: Model Training
# -----
# Logistic Regression
```

```

lr = LogisticRegression()
lr.fit(X_train, y_train)

# Random Forest
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

```

↻

RandomForestClassifier

?

RandomForestClassifier(random_state=42)

◀ ▶

```

# Step 6: Evaluation
# -----
models = {'Logistic Regression': lr, 'Random Forest': rf}

for name, model in models.items():
    y_pred = model.predict(X_test)
    print(f"\n{name} Results")
    print("Accuracy:", accuracy_score(y_test, y_pred))
    print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
    print("Classification Report:\n", classification_report(y_test, y_pred, zero_division=0))

    # ROC AUC (only for binary classification)
    if hasattr(model, "predict_proba") and len(np.unique(y_test)) == 2:
        y_probs = model.predict_proba(X_test)[ :, 1]
        fpr, tpr, _ = roc_curve(y_test, y_probs)
        auc_score = roc_auc_score(y_test, y_probs)
        plt.plot(fpr, tpr, label=f"{name} (AUC = {auc_score:.2f})")

# Plot ROC curve only if applicable
if len(np.unique(y_test)) == 2:
    plt.plot([0, 1], [0, 1], 'k--')
    plt.xlabel("False Positive Rate")
    plt.ylabel("True Positive Rate")
    plt.title("ROC Curves")
    plt.legend()
    plt.show()
else:
    print("ROC curve is not applicable for multiclass classification.")

```

↻

91	1.00	1.00	1.00	1
93	0.00	0.00	0.00	2
97	0.00	0.00	0.00	1
100	0.00	0.00	0.00	1
101	0.00	0.00	0.00	4
106	0.00	0.00	0.00	1
107	0.00	0.00	0.00	1
108	0.00	0.00	0.00	0
111	0.00	0.00	0.00	3
112	0.00	0.00	0.00	1
113	0.00	0.00	0.00	0
114	0.00	0.00	0.00	0
115	1.00	1.00	1.00	1
accuracy			0.30	70
macro avg	0.19	0.20	0.19	70
weighted avg	0.30	0.30	0.30	70

ROC curve is not applicable for multiclass classification

Step 7: Save the Best Model

Save Random Forest model

```
with open("heart_model.pkl", "wb") as f:
    pickle.dump(rf, f)
```

Save encoders

```
with open("label_encoders.pkl", "wb") as f:
    pickle.dump(label_encoders, f)
```

Save scaler

```
with open("scaler.pkl", "wb") as f:
    pickle.dump(scaler, f)
```

```
print("\nModel, encoders, and scaler saved successfully.")
```



Model, encoders, and scaler saved successfully.

```
from sklearn.preprocessing import label_binarize
from sklearn.metrics import auc
```

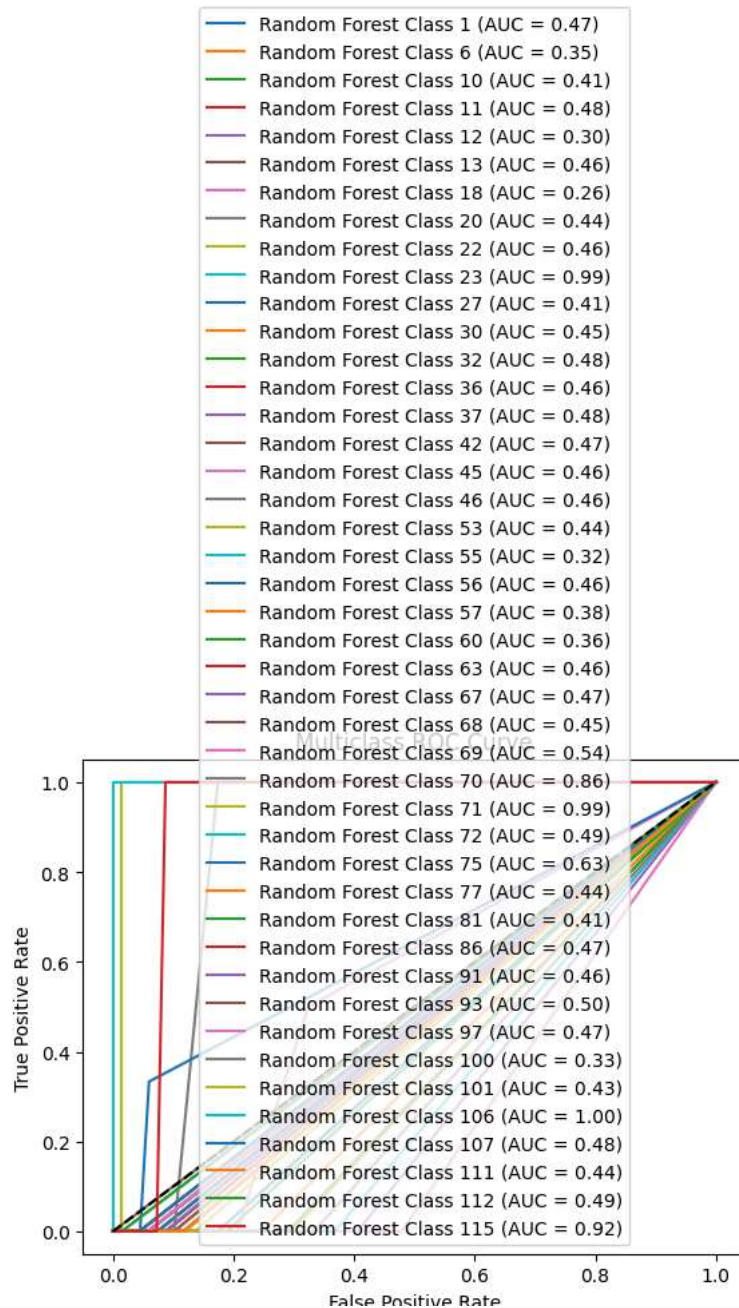
Binarize y_test for multiclass ROC

```
classes = np.unique(y_test)
y_test_bin = label_binarize(y_test, classes=classes)
y_score = model.predict_proba(X_test)
```

Compute ROC curve and AUC for each class

```
for i in range(len(classes)):
    fpr, tpr, _ = roc_curve(y_test_bin[:, i], y_score[:, i])
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, label=f"{name} Class {classes[i]} (AUC = {roc_auc:.2f})")
```

```
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Multiclass ROC Curve")
plt.legend()
plt.show()
```



```
import pandas as pd
df = pd.read_csv('/content/dataset.csv')
df.head()
```



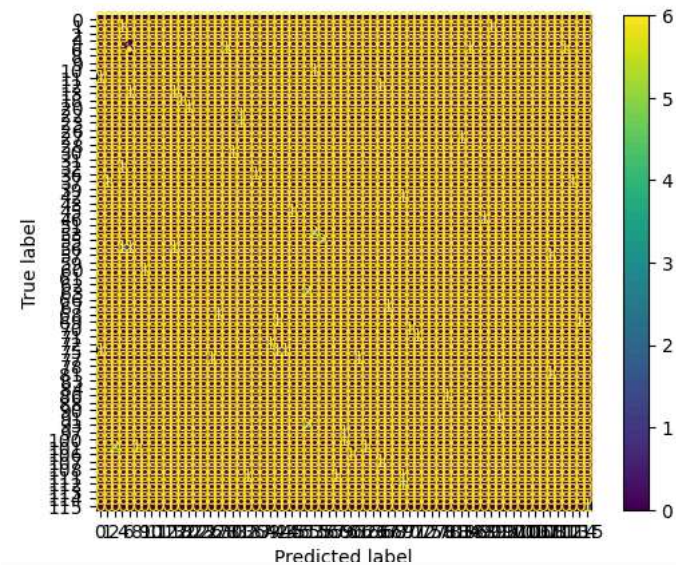
	Disease	Fever	Cough	Fatigue	Difficulty Breathing	Age	Gender	Blood Pressure	Cholesterol Level	Outcome Variable
0	Influenza	Yes	No	Yes	Yes	19	Female	Low	Normal	Positive
1	Common Cold	No	Yes	Yes	No	25	Female	Normal	Normal	Negative
2	Eczema	No	Yes	Yes	No	25	Female	Normal	Normal	Negative
3	Asthma	Yes	Yes	No	Yes	25	Male	Normal	Normal	Positive
4	Asthma	Yes	Yes	No	Yes	25	Male	Normal	Normal	Positive

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)


```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
```

```
ConfusionMatrixDisplay.from_estimator(model, X_test, y_test)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7f34a5a27c90>
```



```
import joblib
joblib.dump(model, 'xgboost_model.pkl')
```

```
['xgboost_model.pkl']
```

```
!pip install streamlit
```

```
Collecting streamlit
```

```
Downloading streamlit-1.45.0-py3-none-any.whl.metadata (8.9 kB)
Requirement already satisfied: altair<6,>=4.0 in /usr/local/lib/python3.11/dist-packages (from streamlit) (5.5.0)
Requirement already satisfied: blinker<2,>=1.5.0 in /usr/local/lib/python3.11/dist-packages (from streamlit) (1.9.0)
Requirement already satisfied: cachetools<6,>=4.0 in /usr/local/lib/python3.11/dist-packages (from streamlit) (5.5.2)
Requirement already satisfied: click<9,>=7.0 in /usr/local/lib/python3.11/dist-packages (from streamlit) (8.1.8)
Requirement already satisfied: numpy<3,>=1.23 in /usr/local/lib/python3.11/dist-packages (from streamlit) (2.0.2)
Requirement already satisfied: packaging<25,>=20 in /usr/local/lib/python3.11/dist-packages (from streamlit) (24.2)
Requirement already satisfied: pandas<3,>=1.4.0 in /usr/local/lib/python3.11/dist-packages (from streamlit) (2.2.2)
Requirement already satisfied: pillow<12,>=7.1.0 in /usr/local/lib/python3.11/dist-packages (from streamlit) (11.2.1)
Requirement already satisfied: protobuf<7,>=3.20 in /usr/local/lib/python3.11/dist-packages (from streamlit) (5.29.4)
Requirement already satisfied: pyarrow>=7.0 in /usr/local/lib/python3.11/dist-packages (from streamlit) (18.1.0)
Requirement already satisfied: requests<3,>=2.27 in /usr/local/lib/python3.11/dist-packages (from streamlit) (2.32.3)
Requirement already satisfied: tenacity<10,>=8.1.0 in /usr/local/lib/python3.11/dist-packages (from streamlit) (9.1.2)
Requirement already satisfied: toml<2,>=0.10.1 in /usr/local/lib/python3.11/dist-packages (from streamlit) (0.10.2)
Requirement already satisfied: typing-extensions<5,>=4.4.0 in /usr/local/lib/python3.11/dist-packages (from streamlit) (4.13.2)
Collecting watchdog<7,>=2.1.5 (from streamlit)
  Downloading watchdog-6.0.0-py3-none-manylinux2014_x86_64.whl.metadata (44 kB)
  44.3/44.3 kB 1.6 MB/s eta 0:00:00
Requirement already satisfied: gitpython!=3.1.19,<4,>=3.0.7 in /usr/local/lib/python3.11/dist-packages (from streamlit) (3.1.44)
Collecting pydeck<1,>=0.8.0b4 (from streamlit)
  Downloading pydeck-0.9.1-py2.py3-none-any.whl.metadata (4.1 kB)
Requirement already satisfied: tornado<7,>=6.0.3 in /usr/local/lib/python3.11/dist-packages (from streamlit) (6.4.2)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.11/dist-packages (from altair<6,>=4.0->streamlit) (3.1.6)
Requirement already satisfied: jsonschema>=3.0 in /usr/local/lib/python3.11/dist-packages (from altair<6,>=4.0->streamlit) (4.23.0)
Requirement already satisfied: narwhals>=1.14.2 in /usr/local/lib/python3.11/dist-packages (from altair<6,>=4.0->streamlit) (1.37.1)
Requirement already satisfied: gitdb<5,>=4.0.1 in /usr/local/lib/python3.11/dist-packages (from gitpython!=3.1.19,<4,>=3.0.7->streamlit) (4.0.1)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas<3,>=1.4.0->streamlit) (2.9.0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas<3,>=1.4.0->streamlit) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas<3,>=1.4.0->streamlit) (2025.2)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.27->streamlit) (3.4.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.27->streamlit) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.27->streamlit) (2.4.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.27->streamlit) (2025.4.14)
Requirement already satisfied: smmap<6,>=3.0.1 in /usr/local/lib/python3.11/dist-packages (from gitdb<5,>=4.0.1->gitpython!=3.1.19,<4,>=3.0.7->streamlit) (3.0.1)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.11/dist-packages (from Jinja2->altair<6,>=4.0->streamlit) (3.0.2)
Requirement already satisfied: attrs>=22.2.0 in /usr/local/lib/python3.11/dist-packages (from jsonschema>=3.0->altair<6,>=4.0->streamlit) (25.1.0)
Requirement already satisfied: jsonschema-specifications>=2023.03.6 in /usr/local/lib/python3.11/dist-packages (from jsonschema>=3.0->altair<6,>=4.0->streamlit) (2025.12.1)
Requirement already satisfied: referencing>=0.28.4 in /usr/local/lib/python3.11/dist-packages (from jsonschema>=3.0->altair<6,>=4.0->streamlit) (0.36.0)
Requirement already satisfied: rpds-py>=0.7.1 in /usr/local/lib/python3.11/dist-packages (from jsonschema>=3.0->altair<6,>=4.0->streamlit) (0.22.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2->pandas<3,>=1.4.0->streamlit) (1.17.0)
```

```

Downloading streamlit-1.45.0-py3-none-any.whl (9.9 MB)
9.9/9.9 MB 48.4 MB/s eta 0:00:00
Downloading pydeck-0.9.1-py2.py3-none-any.whl (6.9 MB)
6.9/6.9 MB 84.1 MB/s eta 0:00:00
Downloading watchdog-6.0.0-py3-none-manylinux2014_x86_64.whl (79 kB)
79.1/79.1 kB 8.3 MB/s eta 0:00:00
Installing collected packages: watchdog, pydeck, streamlit
Successfully installed pydeck-0.9.1 streamlit-1.45.0 watchdog-6.0.0

```

```

import streamlit as st
import joblib
import numpy as np

```

```
model = joblib.load('xgboost_model.pkl')
```

```
st.title("AI Disease Prediction")
```

```

glucose = st.number_input("Glucose Level")
bmi = st.number_input("BMI")
bp = st.number_input("Blood Pressure")
age = st.slider("Age", 20, 100)

```

```

if st.button("Predict"):
    prediction = model.predict([[glucose, bp, bmi, age]])
    st.success(f"Prediction: {'Positive' if prediction[0]==1 else 'Negative'}")

```

2025-05-10 09:26:43.370 WARNING streamlit.runtime.scriptrunner_utils.script_run_context: Thread 'MainThread': missing ScriptRunContext!
 2025-05-10 09:26:43.554
 Warning: to view this Streamlit app on a browser, run it with the following command:

```

streamlit run /usr/local/lib/python3.11/dist-packages/colab_kernel_launcher.py [ARGUMENTS]
2025-05-10 09:26:43.556 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.559 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.561 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.564 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.566 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.569 Session state does not function when running a script without `streamlit run`
2025-05-10 09:26:43.571 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.574 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.577 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.580 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.581 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.582 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.583 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.585 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.587 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.590 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.592 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.594 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.597 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.599 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.601 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.602 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.604 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.605 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.607 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.608 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.611 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.612 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.613 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.615 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-05-10 09:26:43.616 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.

```

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder # Import LabelEncoder
from xgboost import XGBClassifier
import joblib

# Load data
df = pd.read_csv("/content/dataset.csv") # Change filename if needed

# Define target and features
X = df.drop('Outcome Variable', axis=1) # Assuming 'Outcome Variable' is the target

```

```

y = df['Outcome Variable']

# ----> Encode categorical features to numerical using LabelEncoder
# ----> Create a LabelEncoder instance for each categorical column
label_encoders = {}
for col in X.select_dtypes(include=['object']).columns:
    le = LabelEncoder()
    X[col] = le.fit_transform(X[col])
    label_encoders[col] = le # Store the encoder for later use

# 🔄 Step 1: Encode labels to 0, 1, 2, ...
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
y_encoded = le.fit_transform(y)

# 📊 Get number of classes
num_classes = len(np.unique(y_encoded))

# 🔄 Step 2: Scale features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X) # Now X should only contain numerical features

# ✂ Step 3: Split data (stratified)
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y_encoded, test_size=0.2, random_state=42, stratify=y_encoded
)

# 🚀 Step 4: Train XGBoost for multiclass
model = XGBClassifier(objective='multi:softmax', num_class=num_classes, random_state=42)
model.fit(X_train, y_train)

# 💾 Step 5: Save model and label encoder
joblib.dump(model, 'xgboost_model.pkl')
joblib.dump(le, 'label_encoder.pkl')
# ----> Save the label encoders for the categorical features
joblib.dump(label_encoders, 'feature_label_encoders.pkl')

print("✅ Model trained and saved successfully.")

🔄 ✅ Model trained and saved successfully.

!pip install streamlit
!pip install xgboost
!pip install scikit-learn
!pip install joblib
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from xgboost import XGBClassifier
import joblib
import streamlit as st

# ----> Training and saving the model (Run this cell once) <----

# Load data
df = pd.read_csv("/content/dataset.csv") # Update filename if needed

# Define target and features
X = df.drop('Outcome Variable', axis=1)
y = df['Outcome Variable']

# Encode categorical features to numerical using LabelEncoder
label_encoders = {}
for col in X.select_dtypes(include=['object']).columns:
    le = LabelEncoder()
    X[col] = le.fit_transform(X[col])
    label_encoders[col] = le # Store the encoder for later use

# Encode target variable
le = LabelEncoder()
y_encoded = le.fit_transform(y)
num_classes = len(np.unique(y_encoded))

# Scale features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

```

```

# Split data (stratified)
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y_encoded, test_size=0.2, random_state=42, stratify=y_encoded
)

# Train XGBoost
model = XGBClassifier(objective='multi:softmax', num_class=num_classes, random_state=42)
model.fit(X_train, y_train)

# Save model, encoders, and scaler
joblib.dump(model, 'xgboost_model.pkl')
joblib.dump(le, 'label_encoder.pkl')
joblib.dump(label_encoders, 'feature_label_encoders.pkl')
joblib.dump(scaler, 'scaler.pkl') # Added this to save scaler

print("✅ Model trained and saved successfully.")

# ----> Streamlit App (Run this cell to start the app) <----

# Load trained components
model = joblib.load('xgboost_model.pkl')
target_encoder = joblib.load('label_encoder.pkl')
feature_encoders = joblib.load('feature_label_encoders.pkl')
scaler = joblib.load('scaler.pkl') # Load the saved scaler

# Streamlit UI
st.title("AI-Powered Disease Prediction")
st.markdown("Enter patient details to predict possible disease classification.")

# Example inputs - Update based on your dataset
fever = st.selectbox("Fever", ["Yes", "No"])
cough = st.selectbox("Cough", ["Yes", "No"])
fatigue = st.selectbox("Fatigue", ["Yes", "No"])
breathing = st.selectbox("Difficulty Breathing", ["Yes", "No"])
age = st.slider("Age", 0, 120)
gender = st.selectbox("Gender", ["Male", "Female"])
bp = st.number_input("Blood Pressure", min_value=0)
chol = st.number_input("Cholesterol Level", min_value=0)

# Predict button
if st.button("Predict Disease"):
    input_dict = {
        'Fever': fever,
        'Cough': cough,
        'Fatigue': fatigue,
        'Difficulty Breathing': breathing,
        'Age': age,
        'Gender': gender,
        'Blood Pressure': bp,
        'Cholesterol Level': chol
    }

```